



Red Hat Training and Certification

DO188 - Introduction to Containers with
Podman

Travis Michette

Version 1.2

Table of Contents

Introduction	1
Repositories for this Course	1
Demo Setup/Preparing to Teach	1
1. Introduction and Overview of Containers	3
1.1. Introduction to Containers	3
1.1.1. Describing Containers	3
1.1.2. Comparing Containers to Virtual Machines	7
1.1.3. Development for Containers	8
1.2. Introduction to Kubernetes and OpenShift	10
1.2.1. Kubernetes Overview	10
1.2.2. Red Hat OpenShift Container Platform Overview	10
2. Podman Basics	12
2.1. Creating Containers with Podman	12
2.1.1. Introducing Podman	12
2.1.2. Working with Podman	13
2.1.2.1. Pulling and Displaying Images	14
2.1.2.2. Running and Displaying Containers	14
2.1.2.3. Exposing Containers	15
2.1.2.4. Using Environment Variables	16
2.2. DEMO: Creating Containers with Podman	17
2.3. Container Networking Basics	21
2.3.1. Container Networking Basics	21
2.3.2. Managing Podman Networks	21
2.3.3. Enabling Domain Name Resolution	22
2.3.4. Connecting Containers	22
2.4. DEMO: Container Networking Basics	23
2.5. Accessing Containerized Network Services	24
2.5.1. Port Forwarding	24
2.5.1.1. List Port Mappings	24
2.5.2. Networking in Containers	24
2.6. DEMO: Accessing Containerized Network Services	25
2.7. Accessing Containers	28
2.7.1. Container Transparency	28
2.7.2. Start Processes in Containers	28
2.7.3. Open an Interactive Session in Containers	28
2.7.4. Copy Files in and Out of Containers	28

2.8. DEMO: Accessing Containers	29
2.9. Managing the Container Lifecycle	32
2.9.1. Container Lifecycle	32
2.9.2. Inspect a Container	32
2.9.3. Stop a Container	32
2.9.3.1. Stop a Container Gracefully	32
2.9.3.2. Stop a Container Forcefully	32
2.9.3.3. Pause a Container	32
2.9.4. Restarting a Container	32
2.9.5. Remove a Container	32
2.10. DEMO: Accessing Containers	33
3. Container Images	36
3.1. Container Image Registries	36
3.1.1. Container Registries	36
3.1.2. The Containerfile	36
3.1.3. Red Hat Registry	36
3.1.3.1. Useful Container Images	36
3.1.4. Quay.io	36
3.1.5. Manage Registries with Podman	36
3.1.6. Manage Registry Credentials with Podman	36
3.2. DEMO: Container Image Registries	37
3.3. Managing Images	46
3.3.1. Image Management	46
3.3.1.1. Image Versioning and Tags	46
3.3.1.2. Pulling Images	46
3.3.1.3. Building Images	46
3.3.1.4. Pushing Images	46
3.3.1.5. Inspecting Images	46
3.3.1.6. Image Removal	46
3.4. DEMO: Managing Images	47
4. Custom Container Images	52
4.1. Create Images with Containerfiles	52
4.1.1. Creating Images with Containerfiles	52
4.1.2. Choosing a Base Image	52
4.1.3. Containerfile Instructions	52
4.1.4. Container Image Tags	52
4.2. Build Images With Advance Containerfile Instructions	53
4.2.1. Advanced Containerfile Instructions	53

4.2.2. The ENV Instruction	53
4.2.3. The VOLUME Instruction	53
4.2.4. The ENTRYPOINT and CMD Instructions	53
4.2.5. Multistage Builds	53
4.2.6. Examine Container Data Layers	53
4.2.6.1. Cache Image Layers	53
4.2.6.2. Reduce Image Layers	53
4.3. Building Container Images for Use on OpenShift Container Platform (OCP)	53
4.3.1. OpenShift Container Platform Considerations	53
4.3.2. Container Image Design Considerations	54
4.4. Rootless Podman	57
4.4.1. Container Workload Isolation	57
4.4.2. Analyzing Rootless Containers	57
4.4.2.1. Changing the Container User	57
4.4.2.2. Explaining User Mapping	57
4.4.2.3. Limitations of Rootless Containers	57
5. Persisting Data	58
5.1. Volume Mounting	58
5.1.1. Copy-on-write File System	58
5.1.1.1. Implications of a COW File System	60
5.1.2. Store Data on Host Machine	61
5.1.3. Storing Data with Bind Mounts	62
5.1.4. Storing Data with Volumes	62
5.1.5. Storing Data with a tmpfs Mount	62
5.2. Working with Databases	63
5.2.1. Stateful Database Containers	63
5.2.2. Good Practices for Database Containers	63
5.2.3. Importing Database Data	63
5.2.3.1. Database Containers with Data-loading Features	63
5.2.3.2. Data Loading with a Database Client	63
5.2.4. Red Hat Database Containers	63
6. Troubleshooting Containers	64
6.1. Container Logging and Troubleshooting	64
6.2. Remote Debugging Containers	65
7. Multi-Container Applications with Compose	66
7.1. Compose Overview and Use Cases	66
7.1.1. Orchestrate Containers with Podman Compose	66
7.1.2. Podman Pods	66

7.1.3. The Compose File	66
7.1.3.1. Start and Stop Containers with Podman Compose	66
7.1.4. Networking	66
7.1.5. Volumes	66
7.2. Build Developer Environments with Compose	67
7.2.1. Compose Overview	67
7.2.2. Podman Compose and Podman	67
7.2.3. Multi-container Developer Environments with Compose	67
8. Container Orchestration with OpenShift and Kubernetes	68
8.1. Deploy Applications in OpenShift	68
8.2. Multi-pod Applications	69
Appendix A: EX188 Exam Objectives	70

Introduction

Repositories for this Course

Main Repository

The DO188 Demo repository contains the PDF of the instructor notes, and various pre-built demos for the chapters. Demos in this repository are meant to setup and configure the environment automatically to provide consistent demos from course-to-course. Each of the various chapter directories contains one or more playbooks for demonstration or for setting up and configuring the environment for running demonstrations from the various **Demo Repositories**,

- **DO188 Demo:** https://github.com/tmichett/DO467_Demo
- **DO188 (Private Notes):** https://github.com/tmichett/DO467_Notes

The **DO188** Notes repository contains the Asciidoc code for the PDF as well as where the examples get developed. Both of these repositories are part of Jenkins workflows. The primary workflow task monitors the **NOTES** repository for changes. Upon changes, a new PDF is built. The PDF, along with the **Demos** directory are then promoted to the **DEMO** repository and published to the **main** branch publicly for everyone to access (including students).

Demo Setup/Preparing to Teach

The primary demo uses all playbooks to setup an Organizations, Users, Teams, Projects, Credentials, Roles, Job Templates, and finally a Job Workflow Template for approvals.

1. Create Github directory

```
[student@workstation ~]$ mkdir Github ; cd Github
```

2. Clone Repository

```
[student@workstation Github]$  
Cloning into 'DO188_Demo' ...  
remote: Enumerating objects: 118, done.  
remote: Counting objects: 100% (118/118), done.  
remote: Compressing objects: 100% (75/75), done.  
remote: Total 118 (delta 39), reused 88 (delta 12), pack-reused 0  
Receiving objects: 100% (118/118), 734.20 KiB | 5.65 MiB/s, done.  
Resolving deltas: 100% (39/39), done.
```

3. Change to **DO188_Demo** Directory and the Demo Setup directory

```
[student@workstation Github]$ cd D0188_Demo/Demos
```

QUAY Registry Password

There is a Quay repository installed on the **Registry VM** and the credentials are listed below.



<https://registry.ocp4.example.com:8443/repository/>

- **Username:** developer
- **Password:** developer

OCP Credentials and Passwords

<https://console-openshift-console.apps.ocp4.example.com>

Administrator



- **Username:** admin
- **Password:** redhat

Developer

- **Username:** developer
- **Password:** developer

1. Introduction and Overview of Containers

1.1. Introduction to Containers

1.1.1. Describing Containers

- **Container:** A process running from an image containing all required runtime dependencies and components and is self-contained and independent of host operating system and libraries.

A container engine creates a union filesystem which merges the layers of an immutable container image adding a thin readable/writable layer of ephemeral storage which is removed when the container is removed.


Containers use the following Linux components:

- **CGroups** - Partitioning process provided by linux kernel allowing limiting of resources
- **Kernel Namespaces** - Location provided by Linux kernel isolating specific resources from being visible to the entire system and processes. Namespace resources include: network interfaces, process IDs, mount points, IPC resources, and system hostname resources.
- **Secomp**: Limits how processes use system calls and provides a way to whitelist system calls.
- **SELinux**: Security Enhanced Linux providing mandatory access controls for processes. SELinux protects processes from each other and ensures these processes run in a confined space. SELinux has booleans controlling system-level items, File-Contexts, Network Contexts, and Process-Contexts that all work in conjunction to protect the system.


JULIA EVANS
@b0rk

containers = processes

A container is a group of Linux processes



a container process can have 2 PIDs
(or more!)



I started 'top' in a Docker container.
Here's what that looks like in ps:

outside the container


```
$ ps aux | grep top
USER      PID START   COMMAND
root    23540 20:55   top
bork   23546 20:57   top
```

inside the container

```
$ ps aux | grep top
USER      PID START   COMMAND
root      25 20:55   top
```

these two are the same process!

container processes can do anything a normal process can



but you can set rules about what they can do




Figure 1. Containers are Processes - image from Julia Evans (@b0rk)

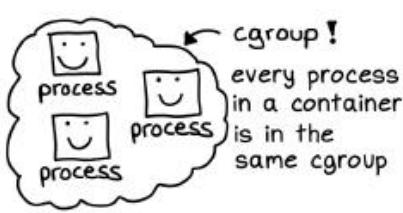
cgroups

13


processes can use a lot of memory



a cgroup is a group of processes



cgroups have memory/CPU {limits}

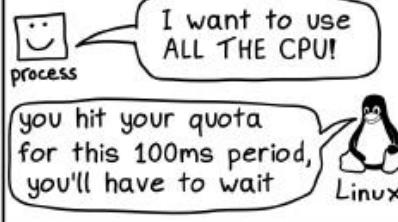


use too much memory: get OOM killed

"out of memory"



use too much CPU: get slowed down



cgroups track memory & CPU usage




Figure 2. CGroup Overview - image from Julia Evans (@b0rk)



Containers aren't magic

These 15 lines of bash will start a container running the fish shell. Try it!
(download this script at bit.ly/containers-arent-magic)

```
wget bit.ly/fish-container -O fish.tar          # 1. download the image
mkdir container-root; cd container-root         #
tar -xf ../fish.tar                            # 2. unpack image into a directory
cgroup_id=$(shuf -i 1000-2000 -n 1)" # 3. generate random cgroup name
cgcreate -g "cpu,cpuacct,memory:$cgroup_id" # 4. make a cgroup &
cgset -r cpu.shares=512 "$cgroup_id"          #     set CPU/memory limits
cgset -r memory.limit_in_bytes=1000000000 \   #
    "$cgroup_id"                                #
cgexec -g "cpu,cpuacct,memory:$cgroup_id" \  # 5. use the cgroup
    unshare -fmuipn --mount-proc \             # 6. make + use some namespaces
    chroot "$PWD" \                           # 7. change root directory
    /bin/sh -c "
        /bin/mount -t proc proc /proc &&      # 8. use the right /proc
        hostname container-fun-times &&       # 9. change the hostname
        /usr/bin/fish"                         # 10. finally, start fish!
```

Figure 3. Containers aren't Magic - image from Julia Evans (@b0rk)

Today, containers have evolved to use container images based on the Open Container Initiative (OCI) standards and specifications (<https://opencontainers.org/>).

Demo of Containers Aren't Magic

Here it is possible to perform a demo of the containers aren't magic to show how the Linux components work that make up containers without using Podman or Docker.

The files and script are available here: https://github.com/tmichett/OCP_Demos/tree/main/Containers_Arent_Magic

This is an excellent demo for courses with Admins, YMMV with a developer course.

Installing CGroup Tools

The default environment doesn't have CGroup tools installed so the **Containers aren't Magic** demo will fail. You MUST install CGroup tools prior to running the demo.



Listing 1. Installing CGroup Tools

```
[root@workstation ~]# yum install libcgroup-tools
```



Also, based on the system setup, in order to properly create the CGroups and run the items, this must be done as a privileged user on the system.

Listing 2. Running the Container Demo as a Privileged User

```
[student@workstation Containers_Arent_Magic]$ sudo  
./containers-arent-magic.sh
```

Images versus Instances

Generally speaking, containers are divided into two components:

- **Container Images:** Image containing the data needed to run a container including all libraries and application components for the container.
- **Container Instance:** A running container based on a container image.

Containers

Generally speaking, when someone refers to a container, it is a **container instance**. A single container image can be used to create multiple **container instances**. Each container instance has the container image data which is **immutable** and one additional unique read/write layer added to the container which is ephemeral and will be removed from the system when the container is removed.



A running container is always based on a container image which is a "read-only" image and filesystem. The container runtime will add an overlay filesystem to the container when it starts up to provide a

temporary read/write layer to the container.



Figure 4. Overlay Filesystem - image from Julia Evans (@b0rk)

1.1.2. Comparing Containers to Virtual Machines

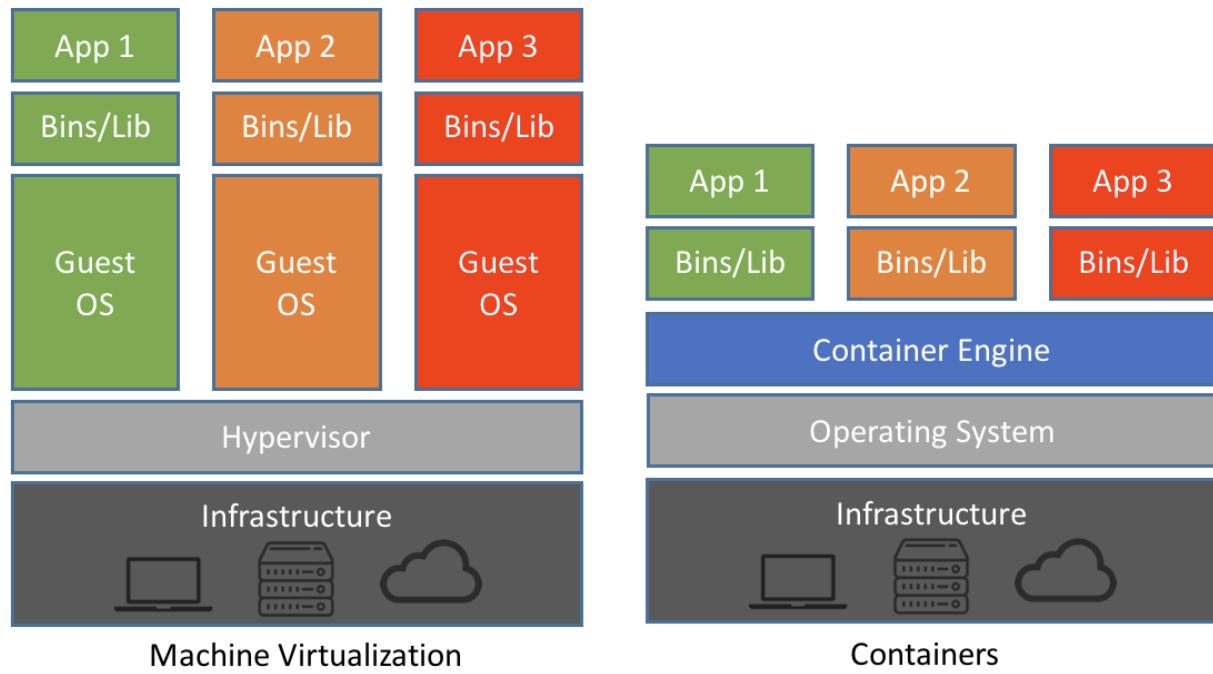


Figure 5. VM Container Comparison

The largest difference between containers and VMs ...

- Image size for VM image is measured typically in Gigabytes because OS and all components are installed, while container images are typically in Megabytes as they only contain the application and libraries needed to run the application.
- VMs require a hypervisor and generally require the same hypervisor, whereas container images can create new containers on any OCI-compliant container engine (providing an OCI compliant container image is used)

Deployment at Scale

Both VMs and Containers can be deployed at scale because they are both based on image templates. The difference is that containers require fewer resources than a virtual machine as VMs require complete operating systems and a hypervisor to be launched whereas containers only require a container engine to be launched making them more portable. Typically, containerized applications are more portable and have larger performance benefits when deployed at enterprise scale.

1.1.3. Development for Containers

Typically, containers are used as a solution for developers as they are easy to use for testing and deployment. Furthermore, based on container images being portable, it is easy to develop a container image and share the image knowing it will work using a supported container engine.

Podman

Red Hat Enterprise Linux uses **podman** to run containers and manage containers with the container engine. Podman will be discussed and used throughout this course and is the replacement for Docker

in Red Hat Enterprise Linux. Podman is meant to manage **pods** which can contain one or more containers. As part of this course, you will be using **podman** to manage individual containers directly, but it is important to note that at enterprise levels with Kubernetes and OpenShift, the smallest managed component is a **pod** which consists of one or more containers.




Figure 6. **podman** Pod Architecture

1.2. Introduction to Kubernetes and OpenShift

1.2.1. Kubernetes Overview

Kubernetes is an orchestration service meant to simplify container management. The smallest unit managed by Kubernetes is a **pod**.

Kubernetes Features

- **Service Discovery and Load Balancing**
- **Horizontal Scaling**
- **Self-Healing**
- **Automated Rollout**
- **Secrets and Configuration Management**
- **Operators**

1.2.2. Red Hat OpenShift Container Platform Overview

OpenShift is an extension of Kubernetes and built on top of Kubernetes. Red Hat OCP adds additional capabilities to Kubernetes which extend and enhance management for the enterprise.



Kubernetes to OCP as Kernel to Red Hat Enterprise Linux Comparisons

In a way, it can be conceptually thought of as if Kubernetes is the Linux Kernel, then Red Hat OpenShift would be equivalent to Red Hat Enterprise Linux. Whereas, Kubernetes is the "Kernel", Red Hat OpenShift Container Platform (OCP) is the "Distribution".

Red Hat OCP Features

- **Developer Workflow**
- **Routes**
- **Metrics and Logging**
- **Unified UI**

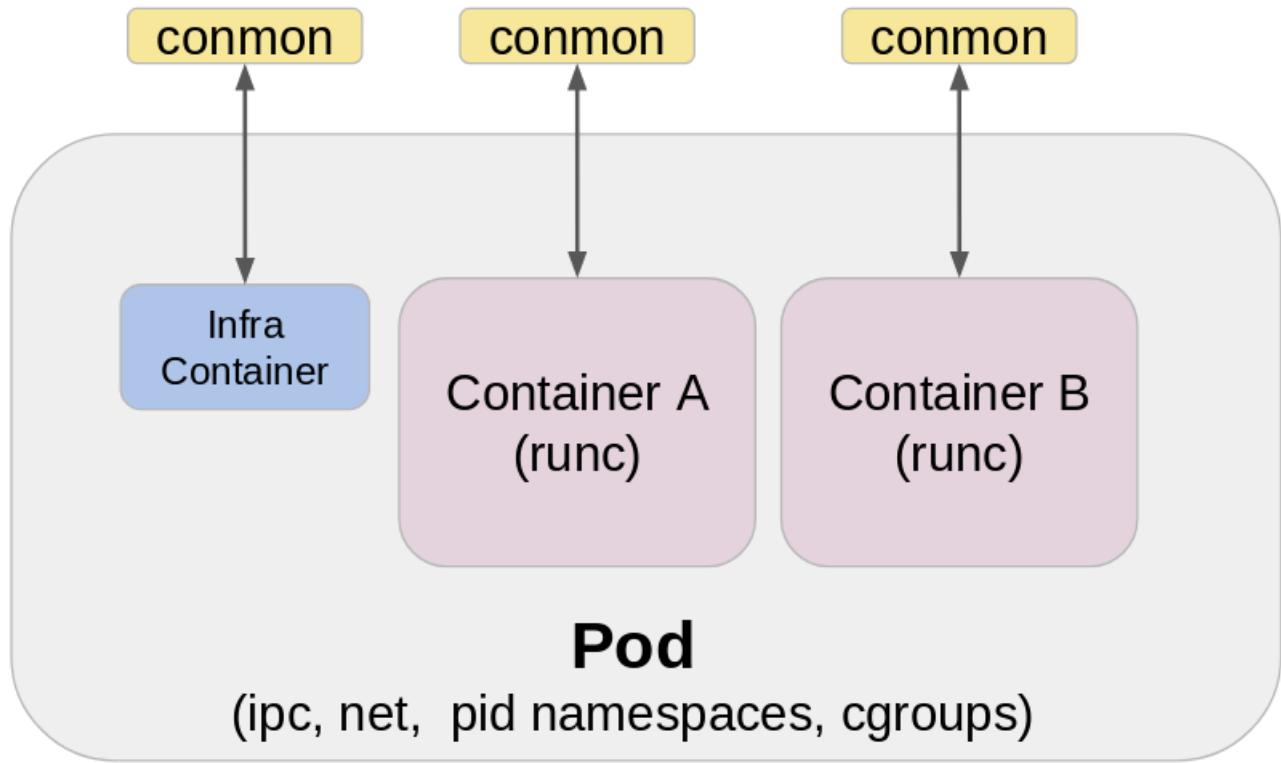


Figure 7. Podman pod Architecture

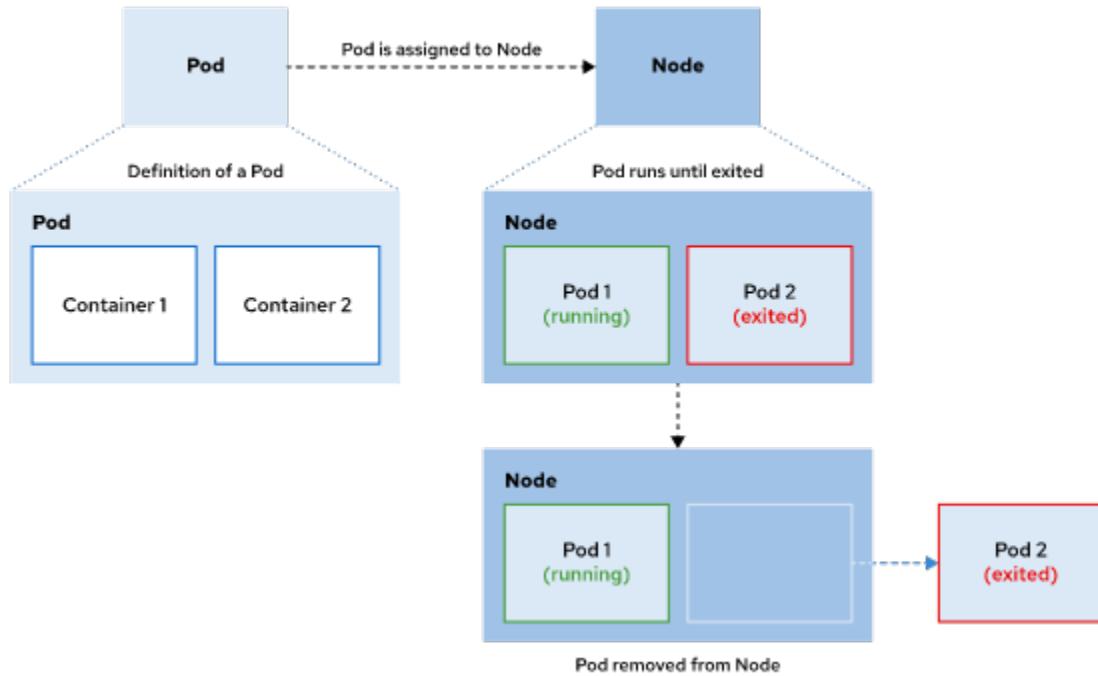


Figure 8. OCP Application and Container Lifecycle

2. Podman Basics

2.1. Creating Containers with Podman

2.1.1. Introducing Podman

Podman is the replacement for Docker and is used to manage containers locally. It allows running, building and deploying containers based on OCI container images. Unlike Docker, **podman** is daemonless. That means that **podman** containers can run without using a privileged daemon and this helps eliminate single-points-of-failure. Podman has a CLI as well as a REST-API.

Almost all **podman** commands are the same as Docker commands, so it is possible to create an alias on the system **alias docker=podman**. It should be understood that **podman** is the CLI management utility and it still relies on an underlying container runtime (**runc**, **crun**, **runv**, and others) which are OCI compliant container runtime environments. Podman containers can be either **root** containers or **non-root/rootless** containers.



podman Command Reference

Podman Commands: <https://docs.podman.io/en/latest/Commands.html>

podman on Red Hat

Podman uses **runc** as the OCI-compliant container runtime on RHEL-based systems.

Container Runtime References:



- **The tool that really runs your containers: deep dive into runc and OCI specifications:** <https://mkdev.me/posts/the-tool-that-really-runs-your-containers-deep-dive-into-runc-and-oci-specifications>
- **Using runc to explore the OCI Runtime Specification:** <https://frasertweedale.github.io/blog-redhat/posts/2021-05-27-oci-runtime-spec-runc.html>
- **Journey From Containerization To Orchestration And Beyond:** <https://iximiuz.com/en/posts/journey-from-containerization-to-orchestration-and-beyond/>

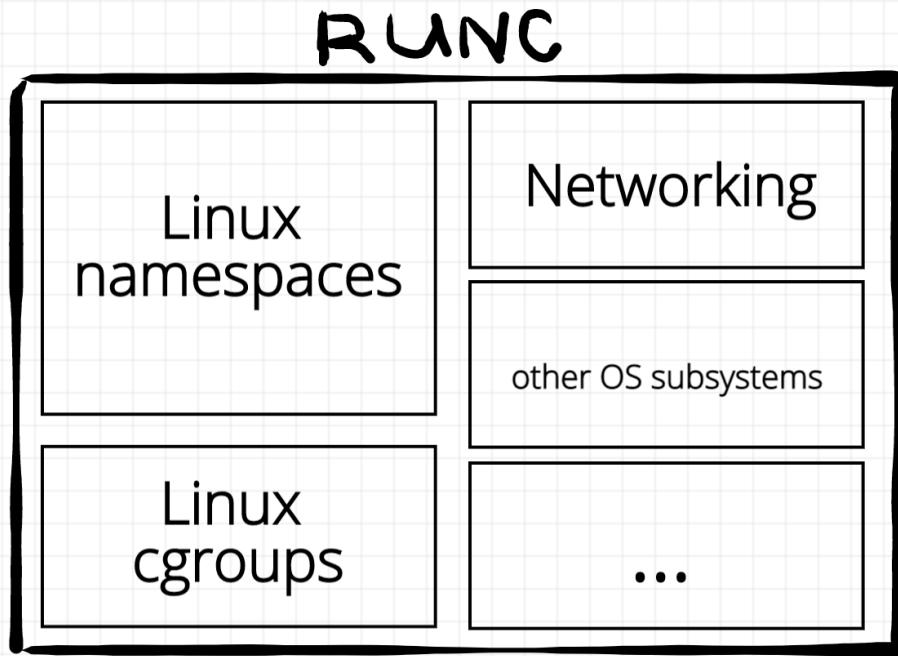


Figure 9. **runc** Container Runtime

2.1.2. Working with Podman

The **podman** package is installed with **Container Tools**, so it is possible to check the version in multiple ways.

Listing 3. RPM Version

```
[student@workstation ~]$ rpm -qa | grep -i podman
podman-catatonit-4.0.2-6.module+el8.6.0+14673+621cb8be.x86_64
podman-4.0.2-6.module+el8.6.0+14673+621cb8be.x86_64
podman-plugins-4.0.2-6.module+el8.6.0+14673+621cb8be.x86_64
podman-gvproxy-4.0.2-6.module+el8.6.0+14673+621cb8be.x86_64
```

*Listing 4. Using the **podman** Command*

```
[student@workstation ~]$ podman -v
podman version 4.0.2
```

*Table 1. **podman** Command Options Covered*

Command Option	Purpose
run	Run and launch a container based on a given image

Command Option	Purpose
<code>run --name <Given_Name></code>	Run and launch a container based on a given image, but name the container based on the <code><Given_Name></code>
<code>run --rm</code>	Run and launch a container based on a given image, but removes the container when exited
<code>ps</code>	List running/active containers
<code>ps -a</code> or <code>ps --all</code>	List all containers regardless of state

2.1.2.1. Pulling and Displaying Images

We learned earlier, containers must have a container image to run. The **podman pull** command will download a container image.

```
[student@workstation ~]$ podman pull quay.io/tmichett/ibm-demo-httd:v1.0
```

Pulling an image will store it locally on the system and can be viewed with **podman images** command.

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
localhost/git_demo_demo	latest	2c7ed9b2d91a	8 days ago	249 MB
quay.io/tmichett/ibm-demo-httd	v1.0	c962db9155dd	2 months ago	258 MB

REMINDER: *Differences between containers and container images*

A container image contains the application and all dependencies for that application to run. A container image is required for a container to run and a container cannot exist without a container image. A container, on the other hand, is an isolated process (runtime environment) based on a container image. The isolation for a container ensures it doesn't disrupt other containers or system processes.



2.1.2.2. Running and Displaying Containers

The **podman run** command is used to launch a container from an image. At a minimum, it will require the container image being launched.

```
[student@workstation ~]$ podman run registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5
echo "I'm the Demo from the container"

I'm the Demo from the container ①
```

① The output from the container is: **I'm the Demo from the container**

The **podman ps** command gives status of all running containers while the **podman ps -a** command gives status of all containers.

```
[student@workstation ~]$ podman ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES

[student@workstation ~]$ podman ps -a
CONTAINER ID IMAGE COMMAND
CREATED STATUS PORTS NAMES
d901895f9102 registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 echo I'm the Demo...
2 minutes ago Exited (0) 2 minutes ago crazy_kalam
```

As container management can get tricky at times for cleanup, it can be a good tip to use the **--rm** option which will remove a container when it exits, as the **--rm** is meant to remove. It is also possible to combine the **--name** option with the run command to give the container a meaningful name.

```
[student@workstation ~]$ podman run --rm --name Demo_Container
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 echo "I'm the Demo from the
container"
I'm the Demo from the container

[student@workstation ~]$ podman ps -a
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
```

podman Output Formats



The output displayed in the screen can be formatted as the **Go** template which is standard, however, it is also possible to use the **--format=format** with most **podman** commands to display the output in JSON format to make things easier to query with something like **jq**.

2.1.2.3. Exposing Containers

The **podman run** command also has additional options which will require the container to expose or allow network traffic as well as running the container in the background like a service daemon.

Table 2. podman run Additional Options

Podman Run Option	Purpose
-p <Local_Port>:<Container_Port>	Exposes and forwards the local port network traffic to the specified container port
-d	Runs the container in the background similar to a daemon.

The commands in the above table will be covered in more detail, but they can be combined to launch things like a webserver or other network workloads.

```
[student@workstation ~]$ podman run -d -p 8080:80 --name travis-demo  
quay.io/tmichett/httpd-custom-demo-new:v1.0
```

```
[student@workstation ~]$ curl localhost:8080  
I am custom material for the D0180 course
```

2.1.2.4. Using Environment Variables

Lastly, we discuss the use of passing environment variables to containers. This allows a way to provide information into a container without placing it as part of the code. This can be useful when using container images like databases or something else where you have a generic image and you customize the container during launch.

Using **podman run -e <Var_Name>=<Value>**, it is possible to pass values for **<Var_Name>**.

2.2. DEMO: Creating Containers with Podman

Registry Requires Authentication

Listing 5. podman login



```
[student@workstation ~]$ podman login registry.ocp4.example.com:8443 -u developer -p developer
Login Succeeded!
```

Example 1. DEMO - Using Podman

1. Explore **podman** to determine versions and downloaded images

```
[student@workstation do188]$ podman -v  
podman version 4.0.2
```

Listing 6. Available Images

```
[student@workstation do188]$ podman images --format "table {{.Repository}}  
{{.Tag}}"  
REPOSITORY TAG  
registry.ocp4.example.com:8443/ubi8/httpd-24 latest  
registry.ocp4.example.com:8443/ubi8/ubi latest  
registry.ocp4.example.com:8443/ubi9/ubi 9.0.0-1468  
registry.ocp4.example.com:8443/ubi8/ubi-minimal 8.5  
registry.ocp4.example.com:8443/redhattraining/hello-world-nginx v1.0
```

2. Use Podman to run a container based on a given image

```
[student@workstation do188]$ podman run registry.ocp4.example.com:8443/ubi8/ubi-  
minimal:8.5 echo "I'm the Demo from the container"  
I'm the Demo from the container
```

Listing 7. Passing Environment Variables to a Container

```
[student@workstation do188]$ podman run -e ENV_VAR1='Demo variable 1' -e ENV_VAR2  
='Demo variable 2' registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 printenv  
ENV_VAR1 ENV_VAR2  
Demo variable 1  
Demo variable 2
```

3. Listing Containers

Listing 8. Running Containers

```
[student@workstation do188]$ podman ps  
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS NAMES
```

Listing 9. All Containers

```
[student@workstation do188]$ podman ps -a
CONTAINER ID  IMAGE                                     COMMAND
CREATED      STATUS          PORTS     NAMES
8916d98e12e5  registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5  echo I'm the
Demo... 10 minutes ago    Exited (0) 10 minutes ago
vibrant_borg
60e09d073a85  registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5  printenv
ENV_VAR1... 9 minutes ago    Exited (0) 9 minutes ago
relaxed_nash
e8aeaa498e47  registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5  echo I'm the
Demo... 2 minutes ago    Exited (0) 2 minutes ago
inspiring_payne
3e5d7231a43d  registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5  printenv
ENV_VAR1... About a minute ago  Exited (0) About a minute ago
friendly_kirch
```

Container Management

It is important to note that unless there is a cleanup or the **podman run** command contains the **--rm** option, when the container finishes a run, it will remain on the system in the "Exited" state. More on this in the **Managing the Container Lifecycle** section.

4. Run a container in the background exposing network traffic (in this case HTTP) traffic to host a website.

```
[student@workstation do188]$ podman run -d -p 8080:80 --name travis-demo
quay.io/tmichett/httpd-custom-demo-new:v1.0
```

Listing 10. Verify Website Data

```
[student@workstation do188]$ curl localhost:8080
I am custom material for the DO180 course
```

5. Accessing **podman** results as JSON

```
[student@workstation do188]$ podman images --format=json | jq '. | map(.Names)'  
[  
  [  
    "registry.ocp4.example.com:8443/ubi8/httpd-24:latest"  
  ],  
  [  
    "registry.ocp4.example.com:8443/ubi8/ubi:latest"  
  ],  
  [  
    "registry.ocp4.example.com:8443/ubi9/ubi:9.0.0-1468"  
  ],  
  [  
    "registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5"  
  ],  
  [  
    "quay.io/tmichett/httpd-custom-demo-new:v1.0"  
  ],  
  [  
    "registry.ocp4.example.com:8443/redhattraining/hello-world-nginx:v1.0"  
  ]  
]
```

**jq**

The **jq** command can be used to filter and process JSON output. In this instance, we wanted the **Names** field, so the **map(.Names)** was used to only return those values.

6. Cleanup Running and stopped containers

```
[student@workstation References]$ podman rm -af
```

2.3. Container Networking Basics

2.3.1. Container Networking Basics

Container networking is setup by default with the **podman** network. By default, all containers are attached to this network and can use it for communicating to each other and the local machine can communicate to the containers through this network. It is possible to use the **podman network** command to create isolated networks for communication between containers.

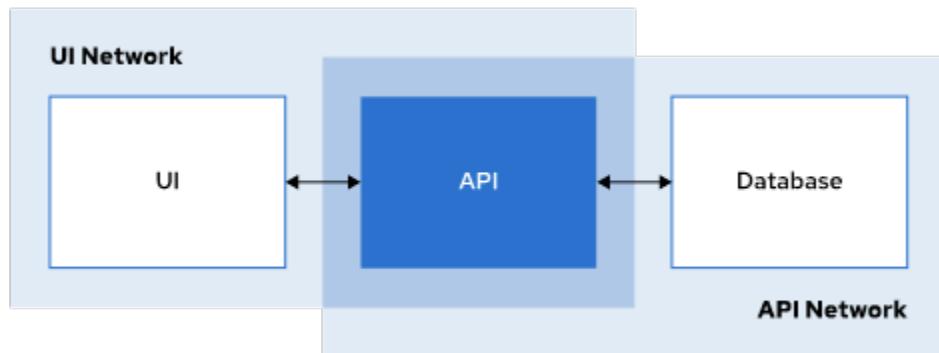


Figure 10. Example of Containers on Different Networks

Default Network Configuration

Earlier versions of **podman** kept network configuration for default networks in `/etc/cni/net.d/87-podman-bridge.conf`. This is no longer true for the default network as **netavark** uses memory-only. Podman default networks do not support DNS resolution either based on backwards compatibility with Docker.



Listing 11. Obtaining Default Network Settings with Podman v4

```
[root@workstation ~]# podman network inspect podman  
[student@workstation ~]$ podman network inspect podman
```

https://github.com/containers/podman/blob/main/docs/tutorials/basic_networking.md

2.3.2. Managing Podman Networks

Table 3. **podman network** Commands

Command	Purpose
podman network create	Creates new podman network

Command	Purpose
podman network ls	Lists podman networks and provides a summary
podman network inspect	Provides detailed JSON output on a network
podman network rm	Removes a network
podman network prune	Removes networks not in use by running containers
podman network connect	Creates a connection between an already running container and an existing container network

2.3.3. Enabling Domain Name Resolution

The default **podman** network has the DNS system disabled. In order to use DNS, a new Podman network must be created and connected to containers. When using DNS the name assigned to the container is the DNS name.



DNS and Default Network

DNS is disabled by default on the **podman** network. It is possible to override the default network and enable DNS, but that is not covered as part of this course.

2.3.4. Connecting Containers

Containers can be connected to one or more Podman networks. Once connected to a network, containers can communicate to other containers on the network via IP address or DNS resolution.

2.4. DEMO: Container Networking Basics

Example 2. DEMO - Using Podman Networks

1. Explore **podman** to determine versions and downloaded images

```
[student@workstation ~]$ podman network create demo-net  
demo-net
```

2. Launch an HTTP Webserver

```
[student@workstation ~]$ podman run -d -p 9080:80 --name travis-demo2 --net demo-  
net quay.io/tmichett/httpd-custom-demo-new:v1.0  
40bfc07475672d28d63b159be2123f771f62c0590327a92f38a91fe72880c646
```

3. Attempt to verify website by DNS name **travis-demo2**

```
[student@workstation ~]$ curl localhost:9080  
I am custom material for the D0180 course  
[student@workstation ~]$ curl travis-demo2  
curl: (6) Could not resolve host: travis-demo2
```

4. Launch RHEL8 Container

```
[student@workstation ~]$ podman run --net demo-net  
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 curl -s travis-demo2  
I am custom material for the D0180 course
```

2.5. Accessing Containerized Network Services

Section Info Here

2.5.1. Port Forwarding

2.5.1.1. List Port Mappings

2.5.2. Networking in Containers

2.6. DEMO: Accessing Containerized Network Services

Example 3. DEMO - Using Podman Port Forwarding

1. Launch a webserver and port forward port **10080** to port **80**

```
[student@workstation ~]$ podman run -d -p 10080:80 --name travis-demo3 --net demo-net quay.io/tmichett/httpd-custom-demo-new:v1.0  
5c422becbd7d4f8e349c2300d98811421d73f4b7b05f3023cdfcccd5631c835d4
```

2. View ports in use using the **podman port** command

```
[student@workstation ~]$ podman port -a  
5c422becbd7d    80/tcp -> 0.0.0.0:10080  
93a63219de88    80/tcp -> 0.0.0.0:9080
```

3. View the IP Address and other Network Settings

```
[student@workstation ~]$ podman inspect travis-demo3 --format=json | jq '.  
|map(.NetworkSettings)|map(.Networks)'
```

4. Retrieve information based on IP address

```
[student@workstation ~]$ podman run --net demo-net  
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 curl -s 10.89.0.7  
I am custom material for the D0180 course
```

5. Demonstrate on default container network

```
[student@workstation ~]$ podman run -d -p 11080:80 --name travis-demo4  
quay.io/tmichett/httpd-custom-demo-new:v1.0
```

Listing 12. Obtaining IP Address

```
[student@workstation ~]$ podman inspect travis-demo4 --format=json | jq '.  
|map(.NetworkSettings)|map(.Networks)'
```

No IP Address



A "rootless" container uses SLIRP Network and doesn't assign an IP Address to the container on the default network namespace. Furthermore, Podman containerized DNS is also not available

6. Demonstrate on default container network as "root" container

```
[student@workstation ~]$ sudo podman run -d -p 12080:80 --name travis-demo5  
quay.io/tmichett/httpd-custom-demo-new:v1.0
```

Listing 13. Obtaining IP Address

```
[student@workstation ~]$ sudo podman inspect travis-demo5 --format=json | jq  
. |map(.NetworkSettings)|map(.Networks)'
```

7. Verify access from workstation machine

```
[student@workstation ~]$ curl 10.88.0.3  
I am custom material for the D0180 course
```



Since we are using the default network from Podman, we have direct access to that network space from the **workstation** machine. Therefore we can access the data within the container by using the container IP address and not relying on port-forwarding.

8. Cleanup all Containers

Listing 14. Cleanup Root Containers

```
[student@workstation ~]$ sudo podman rm -af  
94348119c10655d62bf24565b4b1f0bcde18b7f8f96a9a926e4a5415addf901b
```

Listing 15. Cleanup Rootless Containers

```
[student@workstation ~]$ podman rm -af  
2e27deeb333cd2cd4cc47b3dc1b2bbffce56559224be222f976432dfb3068f1  
3e50448096d5d7caa4a4af9ff5aeb7c6a32b4ded12a2ceb519e35b11225ee2f3
```

2.7. Accessing Containers

2.7.1. Container Transparency

2.7.2. Start Processes in Containers

2.7.3. Open an Interactive Session in Containers

2.7.4. Copy Files in and Out of Containers

2.8. DEMO: Accessing Containers

Example 4. DEMO - Using exec and cp to Access Containers

1. Run the HTTPD container

```
podman run -d -p 2080:80 --name demo-web quay.io/tmichett/httpd-custom-demo-new:v1.0
```

2. View website content

```
[student@workstation ~]$ curl localhost:2080
I am custom material for the D0180 course
```

3. Modify content for website

Listing 16. View content

```
[student@workstation ~]$ podman exec demo-web cat /var/www/html/index.html
I am custom material for the D0180 course
```

Listing 17. Interactively Edit the Content

```
[student@workstation ~]$ podman exec -it demo-web /bin/bash
bash-4.4#
bash-4.4# echo "This is a demo for D0188" > /var/www/html/index.html
bash-4.4# curl localhost
This is a demo for D0188
```

4. View modified website content from workstation.

```
[student@workstation ~]$ curl localhost:2080
This is a demo for D0188
```

5. Copy new website into container

```
[student@workstation ~]$ podman cp ~/Github/OCP_Demos/D0188_Web/index.html demo-
web:/var/www/html/
[student@workstation ~]$ podman cp ~/Github/OCP_Demos/D0188_Web/penguin3.jpeg
demo-web:/var/www/html/
```

6. View website

```
[student@workstation ~]$ curl localhost:2080
```

Instructor Demo - Open in Firefox (Remotely)

Listing 18. Temporarily Add Port



```
[student@workstation ~]$ sudo firewall-cmd --add-port=2080/tcp  
success
```

Then open in a GUI Web Browser!!

2.9. Managing the Container Lifecycle

Section Info Here

2.9.1. Container Lifecycle

2.9.2. Inspect a Container

2.9.3. Stop a Container

2.9.3.1. Stop a Container Gracefully

2.9.3.2. Stop a Container Forcefully

2.9.3.3. Pause a Container

2.9.4. Restarting a Container

2.9.5. Remove a Container

2.10. DEMO: Accessing Containers

Example 5. DEMO - podman and Container Lifecycles

1. Run a simple container using **podman**

```
[student@workstation ~]$ podman run --name simple-demo -e ENV_VAR1='Demo variable 1' -e ENV_VAR2='Demo variable 2' registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 printenv ENV_VAR1 ENV_VAR2
Demo variable 1
Demo variable 2
```

2. Run a **podman ps** to show containers

```
[student@workstation ~]$ podman ps
CONTAINER ID  IMAGE                                     COMMAND
CREATED      STATUS          PORTS          NAMES
42b0e858b01f  quay.io/tmichett/httpd-custom-demo-new:v1.0 /bin/sh -c /usr/s...  2
hours ago    Up 2 hours ago  0.0.0.0:2080->80/tcp  demo-web
```



Only the HTTPD **demo-web** container is running and displayed

3. Run a **podman ps -a** to show all containers

```
[student@workstation ~]$ podman ps -a --format "[{.Image} { {.Names}} { .Status}]"
quay.io/tmichett/httpd-custom-demo-new:v1.0 demo-web Up 2 hours ago
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 simple-demo Exited (0) 6
minutes ago
```

Using --format with Podman



To return specific results, in addition to JSON, it is possible to format data and pull certain fields from the table.

<https://docs.podman.io/en/latest/markdown/podman-ps.1.html>

4. Run a simple container using **podman run --rm** demonstrating the **--rm** option

```
[student@workstation ~]$  
[student@workstation ~]$ podman run --rm --name simple-demo2 -e ENV_VAR1  
='Container erased' -e ENV_VAR2='after this!!'  
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 printenv ENV_VAR1 ENV_VAR2  
Container erased  
after this!!
```



Automatically Cleaning Up Exited Containers

We have learned that the **--rm** command will remove a stopped container. We also know that it can remove a running container when the **-f** or **--force** is used, so what does the **--rm** do when combined with the **podman run** command?

Simply put, the **--rm** command will attempt to remove the container. While the container is running, the command does nothing, as soon as the container "exits", the **--rm** causes the exited container to be removed.

5. Verify that the **simple-demo2** doesn't exist.

```
[student@workstation ~]$ podman ps -a --format "{{.Image}} {{.Names}} {{.Status}}"  
quay.io/tmichett/httpd-custom-demo-new:v1.0 demo-web Up 3 hours ago  
registry.ocp4.example.com:8443/ubi8/ubi-minimal:8.5 simple-demo Exited (0) About  
an hour ago
```

3. Container Images

3.1. Container Image Registries

3.1.1. Container Registries

3.1.2. The Containerfile

3.1.3. Red Hat Registry

3.1.3.1. Useful Container Images

3.1.4. Quay.io

3.1.5. Manage Registries with Podman

3.1.6. Manage Registry Credentials with Podman

3.2. DEMO: Container Image Registries

Example 6. DEMO - Creating a New Repository and Container Image from Github

1. Login to Quay.io



The screenshot shows the Quay.io interface. At the top, there's a navigation bar with links for 'EXPLORE', 'REPOSITORIES', and 'TUTORIAL'. On the right side, there's a search bar, a '+' button, a notification bell with a red badge, and a user profile for 'tmichett'. Below the navigation is a section titled 'Repositories' with a sub-section header '1 - 25 of 216'. A 'Filter Repositories...' button is also present. To the right of the repository list is a sidebar titled 'Users and Organizations' which lists 'tmichett' and 'redhattraining', along with a '+ Create New Organization' button.

Repository Name	Last Modified	Activity	Star
redhattraining / hello-world-nginx	11/26/2019	High	☆
redhattraining / image-tool	04/29/2021	Medium	☆
redhattraining / jupyter-container	09/23/2022	Low	☆
redhattraining / php-hello-dockerfile	05/23/2019	Medium	☆
redhattraining / wordpress	07/06/2021	Medium	☆
redhattraining / nexus	05/24/2022	Low	☆
redhattraining / ad482-ch05s09-connect-cluster	09/22/2021	Medium	☆
redhattraining / nfs-subdir-external-provisioner	07/30/2021	Low	☆

Figure 11. Quay Repositories

2. Click "+ Create New Repository" and select "New Repository"

- Repository Name: **do188_git_demo**
- Visibility: **public**



The screenshot shows the 'Create New Repository' form. At the top, there's a back arrow labeled 'Repositories' and a title 'Create New Repository'. Below this is a breadcrumb navigation bar showing 'tmichett / do188_git_demo' with an orange arrow pointing to the repository name. The form has fields for 'Repository Description' (with a placeholder 'Click to set repository description') and 'Repository Visibility' (with a 'Public' option selected). There are also sections for 'Push' and 'Pull' repository details.

Figure 12. Repository Name and Visibility

- Select **Link to a Github Repository Push**
- Click **Create Public Repository**

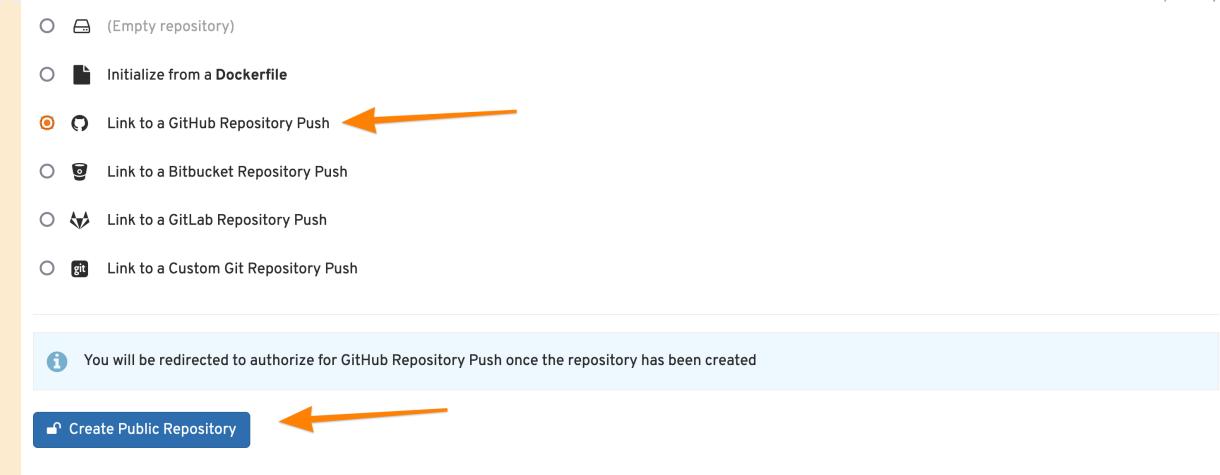


Figure 13. Linking to Version Control Dockerfile/Containerfile

3. Select the Organization containing the Github repository you wish to use and click "Continue"

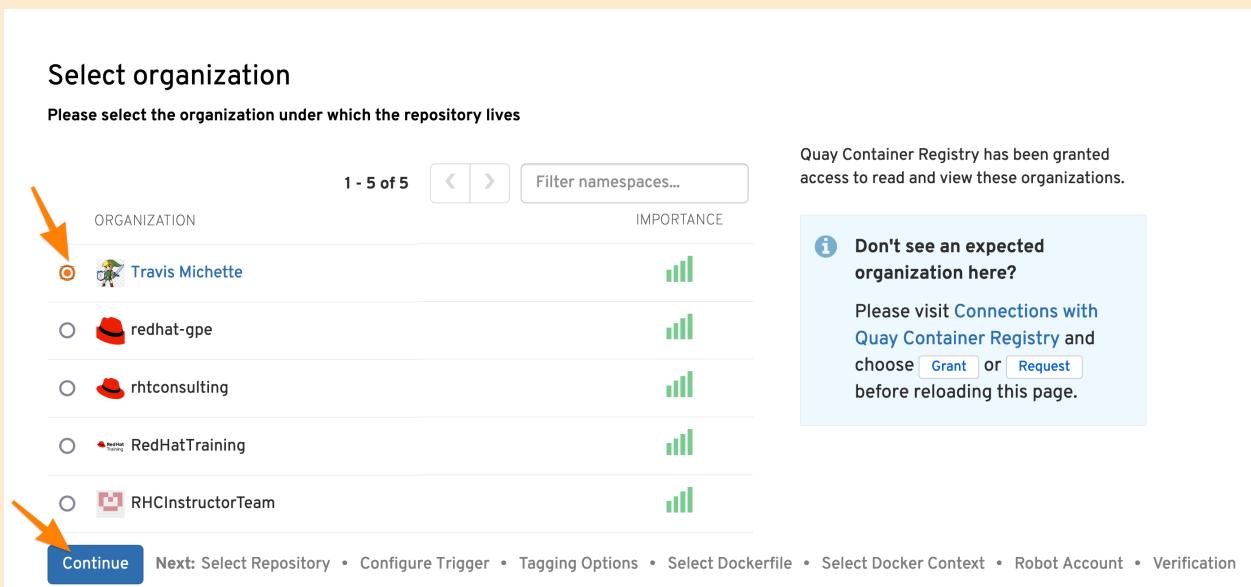


Figure 14. Github Organization Selection

4. Select the Repository you wish to use and click Continue

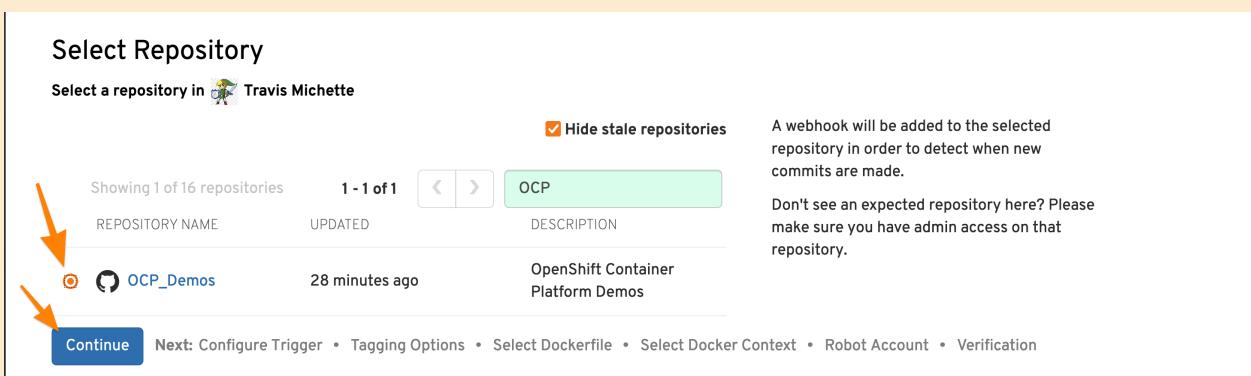


Figure 15. Github Repository Selection

5. Select the automated build trigger options and click Continue

Configure Trigger

Configure trigger options for  Travis Michette/OCP_Demos

Trigger for all branches and tags (default)

Build a container image for each commit across all branches and tags

Do you want to build a new container image for commits across all branches and tags, or limit to a subset?

Trigger only on branches and tags matching a regular expression

Only build container images for a subset of branches and/or tags.

For example, if you use release branches instead of `master` for building versions of your software, you can configure the trigger to only build images for these branches.

All images built will be tagged with the name of the branch or tag whose change invoked the trigger

Continue

Next: Tagging Options • Select Dockerfile • Select Docker Context • Robot Account • Verification

Figure 16. Github Branch and Tag Trigger Option Selection

6. Select the tagging options and click Continue

Configure Tagging

Red Hat E-Business Suite
<https://rsaebs.corp.redhat.com/>

Confirm basic tagging options

Tag manifest with the branch or tag name

Tags the built manifest the name of the branch or tag for the git commit.



By default, all built manifests will be tagged with the name of the branch or tag in which the commit occurred.

Add `latest` tag if on default branch

Tags the built manifest with `latest` if the build occurred on the default branch for the repository.

To modify this default, as well as the default to add the `latest` tag, change the corresponding options on the left.

Add custom tagging templates

No tag templates defined.

Enter a tag template: Add Tag Template

Need more control over how the built manifest is tagged? Add one or more custom tag templates.

For example, if you want all built manifests to be tagged with the commit's short SHA, add a template of `${commit_info.short_sha}`.

As another example, if you want on those manifests committed to a `branch` to be tagged with the branch name, you can add a template of `${parsed_ref.branch}`.

A full reference of for these templates can be found in the [Tag template documentation](#).

Continue

Next: Select Dockerfile • Select Docker Context • Robot Account • Verification

Figure 17. Github Container Tagging Option Selection

7. Select the name and path of the Dockerfile/Containerfile and click Continue

Select Dockerfile

Please select the location of the Dockerfile to build when this trigger is invoked

Continue[Next: Select Docker Context](#) • [Robot Account](#) • [Verification](#)

Please select the location containing the Dockerfile to be built.

The Dockerfile path starts with the context and ends with the path to the Dockerfile that you would like to build

If the Dockerfile is located at the root of the git repository and named Dockerfile, enter `/Dockerfile` as the Dockerfile path.

Figure 18. Github Dockerfile/Containerfile Selection

8. Select the Context directory and click **Continue**

Select Context

Please select the context for the Docker build



/

Continue[Next: Robot Account](#) • [Verification](#)

Please select a Docker context.

The build context directory is the path of the directory containing the Dockerfile and any other files to be made available when the build is triggered.

If the Dockerfile is located at the root of the git repository, enter `/` as the build context directory.

Figure 19. Container Context Build Selection

9. Skip the "Robot Account" and click **Continue**

10. Verify the build trigger and click **Continue**

Ready to go!

Click "Continue" to complete setup of this build trigger.

Continue

Figure 20. Verification of All Options

Authorizing Service

You may be required to authorize the service (Quay) to contact Github. Additionally, when creating the trigger, it will add and updated SSH keys to the source control repository.

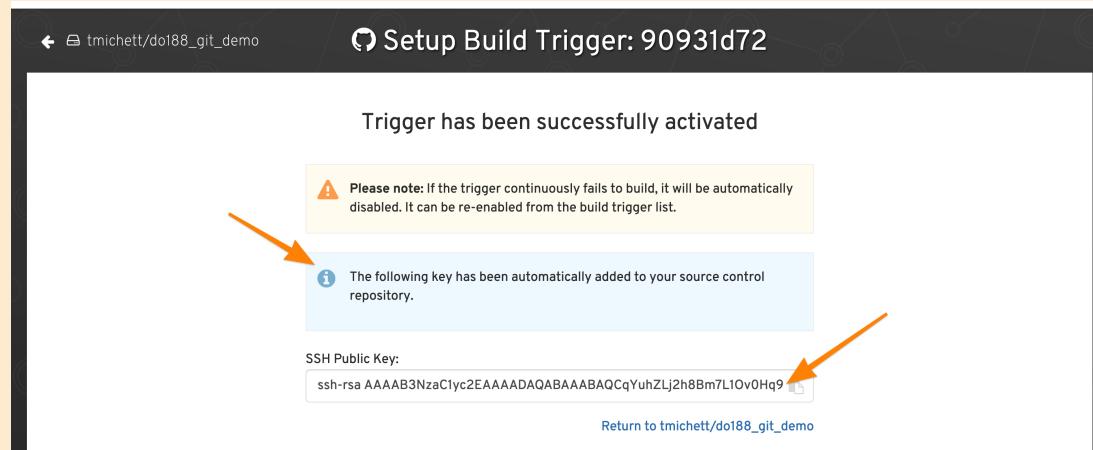
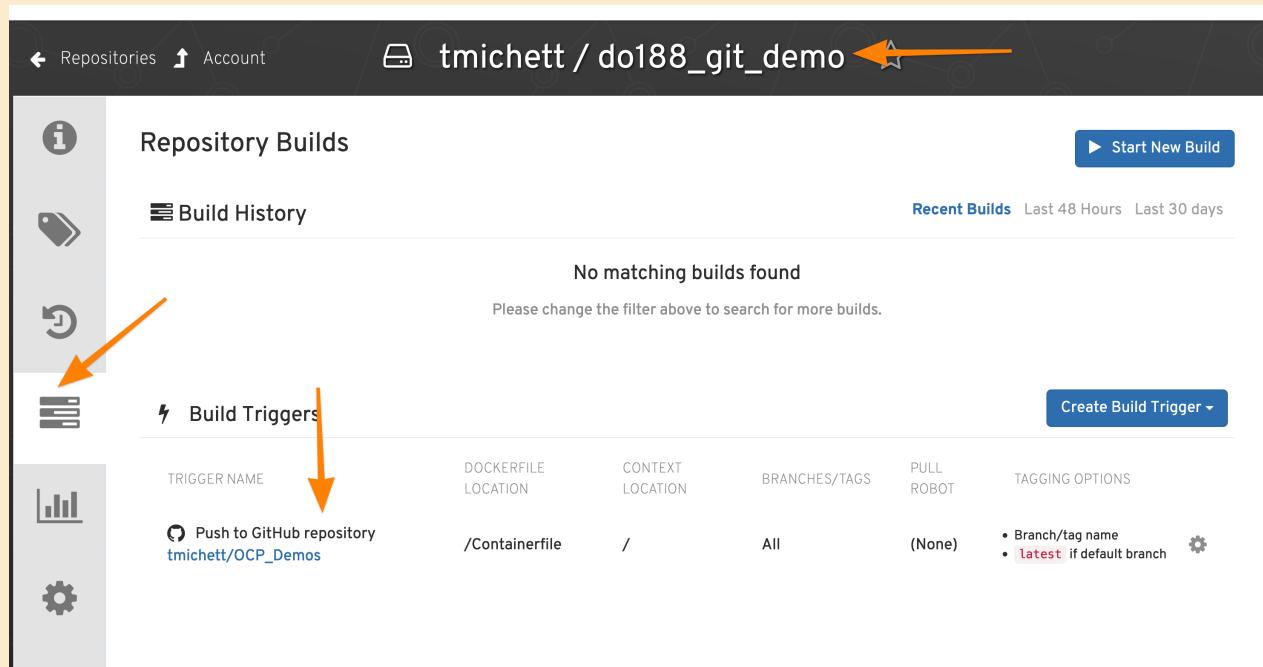


Figure 21. Trigger Activation Confirmation

Testing the Build Trigger

1. Return to the repository and look at the build triggers.



The screenshot shows the "tmichett / do188_git_demo" repository page. On the left sidebar, there are icons for Repository Builds, Build History, Build Triggers (which is highlighted with an orange arrow), and other repository management options. The main content area shows a table for "Build Triggers". One trigger is listed:

TRIGGER NAME	DOCKERFILE LOCATION	CONTEXT LOCATION	BRANCHES/TAGS	PULL ROBOT	TAGGING OPTIONS
Push to GitHub repository tmichett/OCP_Demos	/Containerfile	/	All	(None)	<ul style="list-style-type: none"> Branch/tag name latest if default branch

An orange arrow points to the "Create Build Trigger" button in the top right. Another orange arrow points to the gear icon in the top right corner of the main content area.

Figure 22. Repository Build Triggers

2. Click the Gear icon and select Run Trigger Now

The screenshot shows the GitHub Actions interface for the repository `tmichett/do188_git_demo`. On the left, there's a sidebar with icons for Repository Builds, Build History, and Build Triggers. The main area is titled "Repository Builds" and "Build Triggers". Under "Build Triggers", there is one entry: "Push to GitHub repository `tmichett/OCP_Demos`". To the right of this entry are fields for Dockerfile location (set to "/Containerfile"), context location (set to "/"), branches/tags (set to "All"), and pull robot (set to "(None)"). Below these fields is a "TAGGING OPTIONS" section with two bullet points: "Branch/tag name" and "latest if default branch". A gear icon in this section has a dropdown menu with four options: "View Credentials", "Run Trigger Now" (which is highlighted with an orange arrow), "Disable Trigger", and "Delete Trigger".

Figure 23. Testing Build Trigger

3. On the **Manually Start Build Trigger** leave everything default and click **Start Build**

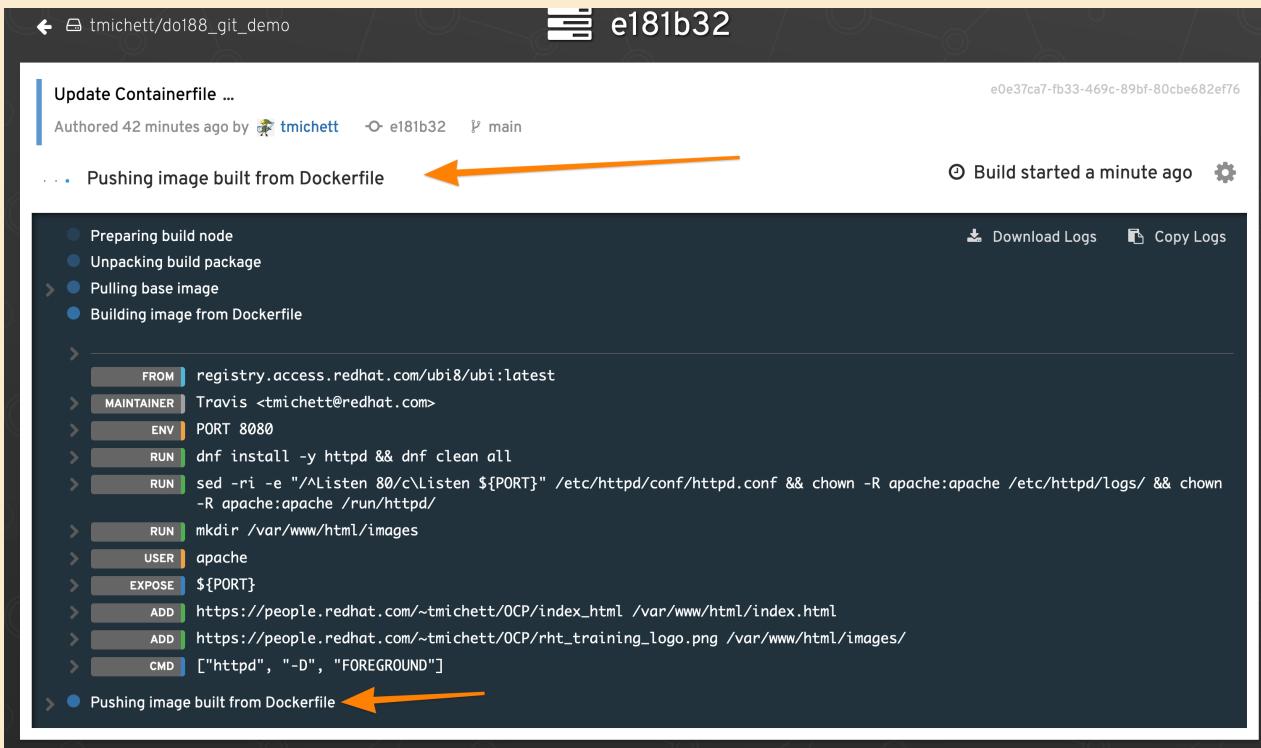
The screenshot shows a modal dialog titled "Manually Start Build Trigger". It contains a section for "Push to GitHub repository `tmichett/OCP_Demos`". Below this is a "Branch/Tag:" input field with a placeholder "Enter or select Branch/Tag". At the bottom right of the dialog are two buttons: "Start Build" (highlighted with an orange arrow) and "Cancel".

Figure 24. Manually Starting Build Trigger

The screenshot shows the "Repository Builds" page for the repository `tmichett/do188_git_demo`. The "Build History" section lists a single build entry: "Update Containerfile ...". This entry has a "BUILD ID" of `e0e37ca7`, was triggered by "Authored 42 minutes ago by tmichett", and was started "Today at 4:39 PM". The build has "latest" and "main" tags. An orange arrow points to the "BUILD ID" column.

Figure 25. Manual Build

4. Click the Build ID and watch the build. Then return to main repository



```

tmichett/do188_git_demo e181b32
Update Containerfile ...
Authored 42 minutes ago by tmichett · e181b32 · main
Pushing image built from Dockerfile →
Build started a minute ago · ⚙️
Preparing build node
Unpacking build package
Pulling base image
Building image from Dockerfile
FROM registry.access.redhat.com/ubi8/ubi:latest
MAINTAINER Travis <tmichett@redhat.com>
ENV PORT 8080
RUN dnf install -y httpd && dnf clean all
RUN sed -ri -e "/^Listen 80/cListen ${PORT}" /etc/httpd/conf/httpd.conf && chown -R apache:apache /etc/httpd/logs/ && chown -R apache:apache /run/httpd/
RUN mkdir /var/www/html/images
USER apache
EXPOSE ${PORT}
ADD https://people.redhat.com/~tmichett/OCP/index.html /var/www/html/index.html
ADD https://people.redhat.com/~tmichett/rht_training_logo.png /var/www/html/images/
CMD ["httpd", "-D", "FOREGROUND"]
Pushing image built from Dockerfile ←

```

Figure 26. Build Log and Verification



Figure 27. Successful Build

5. Attempt to run a container from the new image.

```
[student@workstation ~]$ podman run --rm -d --name do188_custom_ws_demo -p 8080:8080 quay.io/tmichett/do188_git_demo
```

6. Test the website

```
[student@workstation ~]$ curl localhost:8080  
This is a test and demo for the Containerfile build.
```

```

```

3.3. Managing Images

Section Info Here

3.3.1. Image Management

3.3.1.1. Image Versioning and Tags

3.3.1.2. Pulling Images

3.3.1.3. Building Images

3.3.1.4. Pushing Images

3.3.1.5. Inspecting Images

3.3.1.6. Image Removal

3.4. DEMO: Managing Images

Example 7. DEMO - Tagging and Managing Container Images

1. Login and Pull a Container image.

```
[student@workstation do188]$ podman login registry.ocp4.example.com:8443 -u developer -p developer
Login Succeeded!
```

```
[student@workstation do188]$ podman pull
registry.ocp4.example.com:8443/redhattraining/podman-python-server
Trying to pull registry.ocp4.example.com:8443/redhattraining/podman-python-
server:latest...
Getting image source signatures
Copying blob 25ad9a06c050 done

... OUTPUT OMITTED ...

28c663e6c7fd217b7f4d5c9d7d7f21ad8462ad1adc9051ee2bd5cfb80ee34042
```

2. View Images with **podman images**

```
[student@workstation do188]$ podman images
REPOSITORY                                     TAG
IMAGE ID      CREATED      SIZE
quay.io/tmichett/do188_git_demo               latest
2c7ed9b2d91a  4 days ago   249 MB
registry.ocp4.example.com:8443/redhattraining/podman-python-server  latest
28c663e6c7fd  3 months ago  229 MB
```

3. Tag Images

```
[student@workstation do188]$ podman tag 2c7ed9b2d91a git_demo_demo

[student@workstation do188]$ podman tag
registry.ocp4.example.com:8443/redhattraining/podman-python-server
registry.ocp4.example.com:8443/redhattraining/podman-python-server:demo
```

4. View Image Tags

```
[student@workstation do188]$ podman images
REPOSITORY                                     TAG
IMAGE ID      CREATED      SIZE
quay.io/tmichett/do188_git_demo              latest
2c7ed9b2d91a  4 days ago   249 MB
localhost/git_demo_demo                       latest
2c7ed9b2d91a  4 days ago   249 MB
registry.ocp4.example.com:8443/redhattraining/podman-python-server  latest
28c663e6c7fd  3 months ago  229 MB
registry.ocp4.example.com:8443/redhattraining/podman-python-server  demo
28c663e6c7fd  3 months ago  229 MB
```

5. Delete image tag(s) and view images

```
[student@workstation do188]$ podman rmi quay.io/tmichett/do188_git_demo
Untagged: quay.io/tmichett/do188_git_demo:latest
```

```
[student@workstation do188]$ podman images
REPOSITORY                                     TAG
IMAGE ID      CREATED      SIZE
localhost/git_demo_demo                       latest
2c7ed9b2d91a  4 days ago   249 MB
registry.ocp4.example.com:8443/redhattraining/podman-python-server  latest
28c663e6c7fd  3 months ago  229 MB
registry.ocp4.example.com:8443/redhattraining/podman-python-server  demo
28c663e6c7fd  3 months ago  229 MB
```

6. Delete image completely

```
[student@workstation do188]$ podman rmi 28c663e6c7fd
Error: unable to delete image
"28c663e6c7fd217b7f4d5c9d7d7f21ad8462ad1adc9051ee2bd5cfb80ee34042" by ID with more
than one tag ([registry.ocp4.example.com:8443/redhattraining/podman-python-
server:latest registry.ocp4.example.com:8443/redhattraining/podman-python-
server:demo]): please force removal
```

Deleting Images with Multiple Tags

It is possible to remove an image with multiple tags, but you must use the **--force** or **-f** option.

Listing 19. Removing Images with Multiple Tags



```
[student@workstation do188]$ podman rmi 28c663e6c7fd -f
Untagged: registry.ocp4.example.com:8443/redhattraining/podman-
python-server:latest
Untagged: registry.ocp4.example.com:8443/redhattraining/podman-
python-server:demo
Deleted:
28c663e6c7fd217b7f4d5c9d7d7f21ad8462ad1adc9051ee2bd5cfb80ee34042
```

Listing 20. Images and tags fully removed

```
[student@workstation do188]$ podman images
REPOSITORY          TAG      IMAGE ID   CREATED    SIZE
localhost/git_demo_demo latest   2c7ed9b2d91a  4 days ago  249 MB
```



Tags like Hard-Links to inodes

For some people, it is easy to think of images with multiple tags like having hard-links to **inodes**. A tag points to a unique image ID, so tags can be removed without removing the image. When the last tag is removed, the image is removed.

If you remove the image by the ID, it will remove the image and all image tags which is why it is required to use the **-f** or **--force** options. This is similar to deleting an **inode** on a Linux filesystem.

7. Inspecting Images

```
[student@workstation do188]$ podman inspect localhost/git_demo_demo
[
  {
    "Id": "2c7ed9b2d91a13ee511e4d06e04ec602754f151f6b74218437b617a0c7cab1a6",
    "Digest": "sha256:2a4ef2b67a120ee0a10cb2518f8f075b57aa27ff0d13445dbb6a489e2b240baa",
    ...
    ... OUTPUT OMITTED ...
  }
]
```

The podman inspect Command

There are a few things to know about **podman inspect**.



- Images must be downloaded and exist locally on the system in order to be inspected. (*skopeo is used for remote images*).
- Output is extremely verbose, so Go templating can be used with the **--format** option and the curly braces.

8. View environment variables and ports exposed

```
[student@workstation ~]$ podman inspect localhost/git_demo_demo --format
= "{{.Config.Env}}"
[PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin container=oci
PORT=8080]
```

Viewing and Accessing as JSON

It is also possible to use **jq** and process outputs of the commands using JSON filtering and access techniques.



```
[student@workstation ~]$ podman inspect localhost/git_demo_demo | jq '.[].Config.Env'
[
  "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin",
  "container=oci",
  "PORT=8080"
]
```

You can do all kinds of things with **jq** and **grep**.

Listing 21. The jq command and grep Creatively

```
[student@workstation ~]$ podman inspect localhost/git_demo_demo | jq '.[].Config.Env' |
jq -c '.[]' | tostring' | grep PORT
"PORT=8080"
```

4. Custom Container Images

4.1. Create Images with Containerfiles

Section Info Here

4.1.1. Creating Images with Containerfiles

4.1.2. Choosing a Base Image

4.1.3. Containerfile Instructions

DNF vs MicroDNF

It is common, especially for container images, to use MicroDNF as DNF may be generally unavailable based on size and it requiring Python. Traditionally, the **minimal** images will have **microdnf** and DNF will not be available.



- https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux_atomic_host/7/html/getting_started_with_containers_using_red_hat_universal_base_images_standard_minimal_and_runtimes
- <https://fedoramagazine.org/building-smaller-container-images/>

4.1.4. Container Image Tags

4.2. Build Images With Advance Containerfile Instructions

Section Info Here

4.2.1. Advanced Containerfile Instructions

4.2.2. The ENV Instruction

4.2.3. The VOLUME Instruction

4.2.4. The ENTRYPOINT and CMD Instructions

4.2.5. Multistage Builds

4.2.6. Examine Container Data Layers

4.2.6.1. Cache Image Layers

4.2.6.2. Reduce Image Layers

4.3. Building Container Images for Use on OpenShift Container Platform (OCP)

4.3.1. OpenShift Container Platform Considerations

Unlike Podman which runs containers locally based on the container images, OCP is able to leverage container images and the orchestration in the background causes these images to be available to various compute nodes (worker nodes) within the infrastructure and launches containers. When running containers locally, we can easily run containers in a privileged mode using things like "sudo podman run" and performing the bind/volume mounting of containers with the :Z or :z to take care of things like permissions and SELinux constraints.

OpenShift, on the other hand uses multiple different layers of security, RBAC, and other components to control how containers running in addition to all the basic Linux controls (CGroups, Namespaces, SELinux, etc.) that are available. One of these is SCCs (Security Context Constraints) which are beyond the scope of the course. However, this is being mentioned here as sometimes additional "looser" SCCs and Service Accounts may need to be created to allow a container to run from an improperly created Container Image.

Images Needing SCC Changes

Sometimes in order to access items in different containers (based on generic images), it is necessary to change and loosen security restrictions. This would be done for a port and a deployment that has an image for the Apache user, but where permissions within the container image itself aren't setup correctly.



Listing 22. Changing SCC and Permissions

```
$ oc get pod/ocp-demos-59b5744888-8wqxr -o yaml | oc adm policy scc-subject-review -f -
$ oc create sa apache
$ oc adm policy add-scc-to-user anyuid -z apache
$ oc set serviceaccount deployment/ocp-demos apache
```

4.3.2. Container Image Design Considerations

Designing container images for both Podman and OCP start with the same basic constructs and considerations. However, there are additional items of special care when deploying containers from images on a larger container orchestration environment like OCP.

Critical Considerations

- Exposed ports (containers can't bind to protected ports ≤ 1024)
 - Images using ports ≤ 1024 require additional privileges in order to allow binding to those network ports
- Application folders, Data for Applications, and Runtime environment files must be read/writable by the container user.

Listing 23. OCP Compatible Containerfile for supported HTTPD Image

```
FROM registry.access.redhat.com/ubi8/ubi:latest

MAINTAINER Travis <tmichett@redhat.com>

# Environment variable for Apache Port
ENV PORT 8080

RUN dnf install -y httpd && \
    dnf clean all

# Change configuration of Apache WS to listen on the port specified as an environment
# variable in the Containerfile

RUN sed -ri -e "/^Listen 80/c\Listen ${PORT}" /etc/httpd/conf/httpd.conf

RUN mkdir /var/www/html/images

# Demo to simulate new image being built from pipeline
# Copy all files under src/ folder to Apache DocumentRoot (/var/www/html)
ADD https://people.redhat.com/~tmichett/OCP/index_html /var/www/html/index.html
ADD https://people.redhat.com/~tmichett/OCP/rht_training_logo.png /var/www/html/images/

RUN chmod -R +rx /var/www/html ①

## Needed for OpenShift to Set Permissions
# Fix filesystem ownership for writing logs
# Fix filesystem ownership for accessing the runtime PID
# Prevents need from modifying SCC and creating service accounts
RUN chgrp -R 0 /var/log/httpd /var/run/httpd && \
    chmod -R g=u /var/log/httpd /var/run/httpd ②

## Needed for OpenShift to set a generic user
USER 1001 ③

# Start Apache in the foreground
CMD ["httpd", "-D", "BACKGROUND"]
```

① The data files for the HTTPD server based on standard installation are in **/var/www/html** and those must be readable and executable for the web service to share these files as part of a website (this is required for even locally running containers)

② The HTTPD (Apache) service requires access to write to log files (**/var/log/httpd**) and access to the runtime environment, specifically the PID file and location (**/var/run/httpd**). It is necessary to setup the groups and permissions accordingly so that when deployed by OCP, there are no access issues, otherwise there will be errors referring to things like the **HTTPD PID**

③ The container needs to run as a User. In this instance we are specifying a generic and dummy UID.

OpenShift Container Users

By default OCP runs containers using an arbitrary UID, so we want to ensure that the container can run with these same arbitrary UIDs.

Adapting Containerfiles for OCP



- Directories and files that are readable or writeable by processes (inside the container) should be owned by the **root** group and have group READ and WRITE permissions
- Executable files should have group execute permissions
- Processes running in the container cannot listed on privileged ports (<1024)

References



- **Learning container best practices:** https://access.redhat.com/documentation/en-us/openshift_container_platform/4.10/html-single/images/index#creating-images

4.4. Rootless Podman

Section Info Here

4.4.1. Container Workload Isolation

4.4.2. Analyzing Rootless Containers

4.4.2.1. Changing the Container User

4.4.2.2. Explaining User Mapping

4.4.2.3. Limitations of Rootless Containers

5. Persisting Data

5.1. Volume Mounting

5.1.1. Copy-on-write File System

Image Layers



When building container images, each instruction modifying the container filesystem will create a new READ-ONLY data layer for the image. Image layers are stacked as a set of changes from previous layers.

Listing 24. Containerfile

```
FROM registry.access.redhat.com/ubi8/ubi-minimal ①

RUN microdnf install httpd ②
RUN microdnf clean all ③
RUN rm -rf /var/cache/yum ④

CMD httpd -DFOREGROUND ⑤
```

① Base layer of image

② Image layer where application is installed

③ Image layer where DNF cache is cleaned

④ Image layer where the **/var/cache/yum** us removed

⑤ Command launched with container based on this image




Figure 28. Container Image from Containerfile

Each Command Creates an Image Layer

Every command in a Dockerfile/Containerfile creates new layers to the container image (even if contents are removed and not available at runtime). In order to create optimized container images, it is often better practice to have a single instruction creating an image layer.

Listing 25. Optimized Containerfile



```
FROM registry.access.redhat.com/ubi8/ubi-minimal

RUN microdnf install httpd && \
    microdnf clean all && \
    rm -rf /var/cache/yum

CMD httpd -DFOREGROUND
```

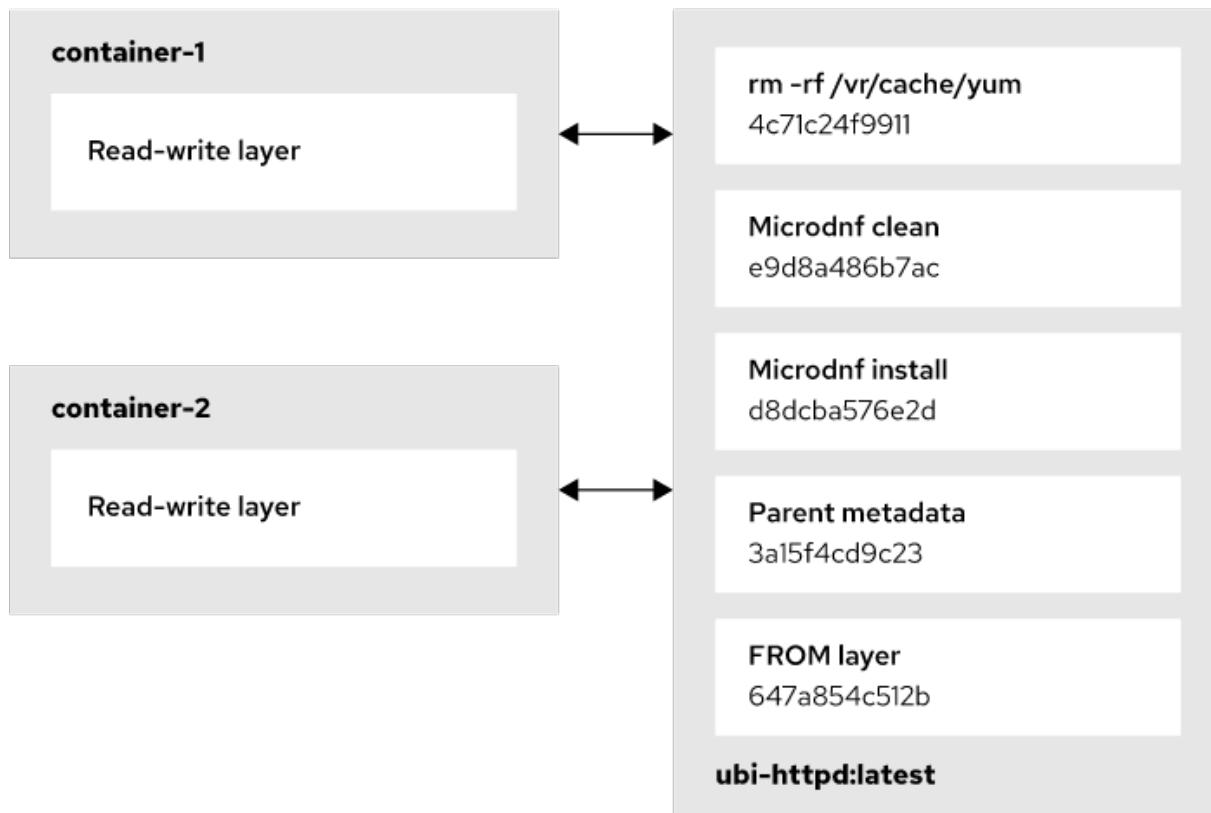


Figure 29. Running Containers - showing Container Image and the Read/Write Ephemeral Layer

Multiple Containers based on Single Container Image

You should recall that there can be multiple containers based on a single container image. Each running container will get its own READ/WRITE filesystem layer that is destroyed when the container is destroyed.



By default, the `/etc/containers/storage.conf` file controls where storage is used. For **root-based** containers, we use an Overlay filesystem. For **rootless** containers, FUSE-Overlay is used. This file also defines the temporary storage location for the READ/WRITE layers on the filesystem.

5.1.1.1. Implications of a COW File System

COW (Copy-on-Write) is the filesystem used for many VMs and containers. This allow the READ/WRITE layer using a union filesystem as shown in the images above.

Image Layers and Sharing

COW and container images are shared between different containers because the container image layers are immutable (read-only) layers



Performance Issues



Based on how union filesystems behave, it is possible to introduce performance issues for write-intensive processes. Because the COW layer for READ/WRITE will be a larger difference between that and the container image.

5.1.2. Store Data on Host Machine

In most instances it is desirable to store data for use in containers. This is called persistent storage of data and is accomplished using external mounts by using either **volumes** or **bind mounts**.

Container Volumes



In previous courses, we focused on **bind mounts** like in the DO180 course. Using these mounts accomplished persistent data storage, but required using things like :Z when mounting or setting SELinux and other filesystem permissions to function properly. Volumes on the other hand are fully managed by Podman and are generally though as easier to use (especially when moving to an orchestration tool like Kubernetes and OCP which use volumes).

- Persistence
- Use of Host File System
- Ease of Sharing

Volumes vs. Bind Mounts

It should be noted that **Volumes** are managed by Podman, while **Bind Mounts** are mounts managed by the user.

The Podman command can mount **volumes** or **bind mounts** using the **-v** or **--volume** parameter.

Listing 26. Typical Mounting options for Volumes

```
--volume /path/on/host:/path/in/container:OPTIONS
```

Listing 27. Typical Mounting options for Mount

```
--mount type=TYPE,source=/path/on/host,destination=/path/in/container
```

The **mount** command is preferred way to mount directories in a container, but the **-v** parameter is the most widely used.

Using podman mount

The **--mount** option specifies and identifies the volume type



Volume Types

- **bind** - Used for Bind mounts
- **volume** - Used for Volume mounts
- **tmpfs** - Used for creating read-only (ephemeral mounts)

5.1.3. Storing Data with Bind Mounts

Bind mounts can mount existing directories on the host filesystem into the container. However, it is important to note that things like SELinux and permissions must also be taken into account. The **:Z** can often be used and assist with this.

5.1.4. Storing Data with Volumes

Volumes allow Podman to manage data and mounts. The **podman volume** command is used to create volumes and leverage volumes.

Importing and Exporting Data in Volumes

The **podman volume** command supports both **import** and **export** arguments which is capable of importing or exporting data into a Podman volume.



By default, the volumes are stored on the host system

- **rootless:** \$HOME/.local/share/containers/storage/volumes/

5.1.5. Storing Data with a tmpfs Mount

5.2. Working with Databases

Section Info Here

5.2.1. Stateful Database Containers

5.2.2. Good Practices for Database Containers

5.2.3. Importing Database Data

5.2.3.1. Database Containers with Data-loading Features

5.2.3.2. Data Loading with a Database Client

5.2.4. Red Hat Database Containers

6. Troubleshooting Containers

6.1. Container Logging and Troubleshooting

Section Info Here

6.2. Remote Debugging Containers

Section Info Here

7. Multi-Container Applications with Compose

7.1. Compose Overview and Use Cases

Section Info Here

7.1.1. Orchestrate Containers with Podman Compose

7.1.2. Podman Pods

7.1.3. The Compose File

7.1.3.1. Start and Stop Containers with Podman Compose

7.1.4. Networking

7.1.5. Volumes

7.2. Build Developer Environments with Compose

7.2.1. Compose Overview

7.2.2. Podman Compose and Podman

7.2.3. Multi-container Developer Environments with Compose

8. Container Orchestration with OpenShift and Kubernetes

8.1. Deploy Applications in OpenShift

Section Info Here

8.2. Multi-pod Applications

Section Info Here

Appendix A: EX188 Exam Objectives

The **EX188** exam objectives are listed here: <https://www.redhat.com/en/services/training/ex188-red-hat-certified-specialist-containers-exam?section=Objectives>. For convenience, a snapshot of the objectives have been placed in this guide.

- Implement images using Podman
 - Understand and use FROM (the concept of a base image) instruction.
 - Understand and use RUN instruction.
 - Understand and use ADD instruction.
 - Understand and use COPY instruction.
 - Understand the difference between ADD and COPY instructions.
 - Understand and use WORKDIR and USER instructions.
 - Understand security-related topics.
 - Understand the differences and applicability of CMD vs. ENTRYPOINT instructions.
 - Understand ENTRYPOINT instruction with param.
 - Understand when and how to expose ports from a Containerfile.
 - Understand and use environment variables inside images.
 - Understand ENV instruction.
 - Understand container volume.
 - Mount a host directory as a data volume.
 - Understand security and permissions requirements related to this approach.
 - Understand the lifecycle and cleanup requirements of this approach.
- Manage Images
 - Understand private registry security.
 - Interact with many different registries.
 - Understand and use image tags
 - Push and pull images from and to registries.
 - Back up an image with its layers and meta data vs. backup a container state.
- Run containers locally using Podman
 - Run containers locally using Podman
 - Get container logs.
 - Listen to container events on the container host.
 - Use Podman inspect.

- Specifying environment parameters.
- Expose public applications.
- Get application logs.
- Inspect running applications.
- Run multi-container applications with Podman
 - Create application stacks
 - Understand container dependencies
 - Working with environment variables
 - Working with secrets
 - Working with volumes
 - Working with configuration
- Troubleshoot containerized applications
 - Understand the description of application resources
 - Get application logs
 - Inspect running applications
 - Connecting to running containers